


Python Memory Management

Friday, January 22, 2016 7:51 PM

- Python structures use substantially more(3-6x) memory than C++ ones
- Inefficient memory allocation, does not remove de-allocated memory of objects
 - 256kb Arena divided into 4kb pools, with fixed size blocks
 - Objects go into various 4kb pools depending on the specified block sizes
 - freepools, usedpools, partially_allocated_arena are some of the variables which keep track
 - Once created, the memory in the pools and arenas is reused when necessary, but the created chunks aren't deleted - this builds up when different pools needed



- Pickling also uses an unnecessarily large chunk of memory
 - More efficient to serialize certain objects into a file
 - To control what is pickled within a class, use(long term storage):
 - `_getstate_(self)`: assigns what is passed(variables and such) to the pickle dump
 - `_setstate_(self)`: assigns what pickle retrieves and variables it sets
 - Can also pass the entire dictionary and remove or set the necessary variables(not very robust)
 - Use `protocol=pickle.HIGHEST_PROTOCOL` as parameter in a pickle dump to get best efficiency(check which protocols are compatible with the appropriate version of Python)
 - Use the `theano.pkl_utils` module to pickle Theano objects in numpy readable form
 - Good for sharing
- Theano memory model
 - Manages its memory apart from Python
 - Shared variables in the buffer are unique(never aliased to another shared variable)
 - `borrow=True` flag allows aliasing of variables(faster, but more prone to bugs)
 - Shared variables by default are not aliasable(`borrow=False`)
 - deep copy of an object is made when passed to shared, otherwise, the shared variable uses the same object that's passed to it
 - Use for large objects
 - Useless with GPUs
 - Also can be used with `get_value()` - both may create copies
 - Returns a value that may be aliased if `borrow=True`
 - Returns a value that may not be aliased if `borrow=False`
 - `return_internal_type=True` returns the actual internal representation unlike `borrow=True` with the GPU(may use a different internal representation)
 - `get_value()` always returns the same object that is created as the shared variable
 - Safe to use `return_internal_type=True` when the value won't be modified
 - Usage with `set_value()` - reuses provided buffer
 - Pattern:

```
s.set_value(  
    ◆ some_inplace_fn(s.get_value(borrow=True)),  
    borrow=True)
```
 - Use contiguous values when assigning to `CudaNdarraySharedVariable`

- Works with defining a function

```
import theano, theano.tensor
```

- ```
x = theano.tensor.matrix()
y = 2 * x
f = theano.function([theano.In(x, borrow=True)], theano.Out(y, borrow=True))
```

  - ◆ theano.In and theano.Out arguments as the arguments and the function values
  - ◆ Reuses the same buffer